

LFQ: Online Learning of Per-flow Queuing Policies using Deep Reinforcement Learning

Maximilian Bachl, Joachim Fabini, Tanja Zseby

Technische Universität Wien, Vienna, Austria

Fair Queuing

- Separate each flow in separate queue
- No flow can “steal” bandwidth
- Popular implementation for Linux: fq


Fair Queuing

- Separate each flow in separate queue
- No flow can “steal” bandwidth
- Popular implementation for Linux: fq


Static buffer size:

Buffer often too large or too small


Buffer too large



Buffer too small



Buffer perfect



Deep Reinforcement Learning for Buffering

- Have a separate queue for each flow in a switch/router (fair queuing)
- Define some reward function (like high throughput, low delay)
- Use **Deep Reinforcement Learning**
 - to **fingerprint** each flow
 - and to use the **correct buffer size** for the flow **based on experience**, which **maximizes the reward**
- We call it *Learning Fair Queuing*

Reward

$$\textit{Reward} = \textit{throughput} - \alpha \times \textit{queue size}$$

Input features

- queue size
- standard deviation of the queue size
- maximum allowed buffer size
- rate of incoming data
- rate of outgoing data
- time since the last packet loss

Neural Network

- It is a regression problem: Predict optimal buffer size for the flow based on the inputs
- Fully connected neural network with three layers and leaky ReLU
- Outputs optimal maximum buffer size each time packet is received
 - To save compute it is possible, for example, only output optimum buffer size every 100 ms

Implementation

- ns-3 for network simulation
- Pytorch's C++ API for Deep Learning
- Integrate Pytorch into ns-3
- Develop queuing discipline based on fair queuing
- Each flow is managed by Deep Reinforcement Learning

Offline Learning

- Randomly sample bandwidth, delay, congestion control and flow duration
- Simulate each flow continuously using the optimal buffer size output by the Deep Reinforcement Learning
- For each flow, at a random time, launch experiment (A/B testing):
 - Continue one experiment with current buffer size +1 packet and one with current buffer size -1 packet (two different experiments)
- Let the neural network learn which buffer size achieved the higher reward
- Train using L1 loss (mean absolute error), using a couple of results together as a batch

Training takes several hundred thousand flows to converge.

Measuring success

Should learn

- Larger bandwidth \rightarrow larger buffer
- Larger delay \rightarrow larger buffer
- New Reno Congestion Control should get bigger buffer than BIC

Correlations

Correlations calculated of each of the above plots:
Correlation of 100% \rightarrow Our Reinforcement Learning system learned successfully

Table: Correlation between bandwidth/delay of the link and output buffer size.

| Congestion control | bandwidth | delay |
|--------------------|-----------|-------|
| New Reno | 98.3% | 99.1% |
| BIC | 80.2% | 90.2% |


Average queue sizes

Average queue is smaller for New Reno: This is expected since it has a larger multiplicative decrease parameter!


→ New Reno needs a larger buffer on average so that the buffer never becomes empty. Our mechanism learned that!

| Congestion control | avg. max. queue length |
|--------------------|------------------------|
| New Reno | 22 |
| BIC | 8 |


25 Mbit/s, 15 ms, New Reno, LFQ



6 Mbit/s, 15 ms, New Reno, LFQ




25 Mbit/s, 15 ms, BIC, LFQ



Average queue sizes

| Congestion control | avg. max. queue length |
|--------------------|------------------------|
| New Reno | 13 |
| BIC | 5 |

6 Mbit/s, 15 ms, New Reno, LFQ



Online Learning

Same as Offline Learning except:

- Don't do A/B testing
- Instead: Use second neural network that outputs expected reward (*Critic Network*) (trained using L2 loss (mean squared error)).
- Either perform experiment A **or** experiment B (either current buffer size -1 packet or +1 packet)
- If result was better than what the Critic Network expected, let first neural network (*Actor Network*) learn to output this buffer size in the future.

Training takes longer (several million training flows)

Correlations

Table: Correlation between bandwidth/delay of the link and output buffer size.


| Congestion control | bandwidth | delay |
|--------------------|-----------|-------|
| New Reno | 86.1% | 87.7% |
| BIC | 72.9% | 73.6% |

Average queue sizes

| Congestion control avg. max. queue length | |
|--|----|
| New Reno | 14 |
| BIC | 8 |


Example: Our solution

15 Mbit/s bandwidth, 5 ms delay




Example: FqCoDel (default on home routers)

15 Mbit/s bandwidth, 5 ms delay



Example: Fifo (default on Linux)

15 Mbit/s bandwidth, 5 ms delay



Systematic comparison

Table: 400 experiments; delay 5 – 25 ms; bandwidth 5 – 25 Mbit/s;
New Reno and Bic Congestion Control. Results averaged.

| | avg. throughp. | queue size | |
|------------------------------|----------------|------------|-------|
| | | max. | avg. |
| LFQ, offline $\alpha = 0.01$ | 13.4 | 23.9 | 7.7 |
| LFQ, offline $\alpha = 10$ | 12.5 | 12.7 | 3.4 |
| LFQ, online $\alpha = 10$ | 12.8 | 16.1 | 4.5 |
| FqCoDel | 13.7 | 155.4 | 15.4 |
| fq 100 | 11.7 | 100 | 51.1 |
| fq 1000 | 11.9 | 1000 | 630.4 |

Observations

- Scaling of features very important
- Pytorch better to integrate with other code (ns-3) than TensorFlow
- We thought simulating a flow in ns-3 would be faster than running it in the real world. It is not.

Conclusion

- LFQ is based on fair queuing
- It fingerprints each flow
- It learns to optimize a reward function
- It achieves high throughput and low delay when compared to competing solutions
- It has low computational overhead
- We envision deployment close to end users (for example on home routers)

LFQ: Online Learning of Per-flow Queuing Policies using Deep Reinforcement Learning

Maximilian Bachl, maximilian.bachl@tuwien.ac.at
Joachim Fabini
Tanja Zseby

Technische Universität Wien, Vienna, Austria