

LFQ: Online Learning of Per-flow Queuing Policies using Deep Reinforcement Learning

Maximilian Bachl, Joachim Fabini, Tanja Zseby

Technische Universität Wien, Vienna, Austria

“Bufferbloat”

- If queues in routers/switches too small: **Underutilization**
- **Solution:** Make queues very **large** for maximum throughput!
- **New problem:** Packets **wait** a long time (several seconds) in the queue: Bufferbloat

CoDel Active Queue Management against Bufferbloat

- Moving time window of 100 ms
- Queuing Delay **must be** $< 5 \text{ ms}$ once in each window
- Otherwise: Drop packet(s)

Fair Queuing

- Separate each flow in separate queue
- No flow can “steal” bandwidth
- Popular implementation for Linux: fq

Fair Queuing

- Separate each flow in separate queue
- No flow can “steal” bandwidth
- Popular implementation for Linux: fq

Static buffer size:

Buffer often too large or too small

`fq_codel` Combining fair queuing with CoDel

- State-of-the-art
- Separate queues
- Each queue managed by CoDel
- Keeps each flow's queue < 5 ms
- Linux implementation: `fq_codel`


`fq_codel` Combining fair queuing with CoDel

- State-of-the-art
- Separate queues
- Each queue managed by CoDel
- Keeps each flow's queue < 5 ms
- Linux implementation: `fq_codel`


Inadequate interaction with Congestion Control:

Cubic doesn't achieve full throughput or keeps standing queue


Buffer too large



Buffer too small



Buffer perfect



cocoa qdisc¹

- Fair queuing-based qdisc which **adapts** buffer **depending on congestion control**
- Manages to achieve optimal throughput while keeping **delay** from buffering **minimal**
- **Works** for common congestion controls

¹*Cocoa: Congestion Control Aware Queuing*, Bachl, Fabini, Zseby, 2020

cocoa qdisc¹

- Fair queuing-based qdisc which **adapts** buffer **depending on congestion control**
- Manages to achieve optimal throughput while keeping **delay** from buffering **minimal**
- **Works** for common congestion controls

Hand-crafted algorithm:

Works well for current congestion controls but might fail badly for new congestion controls.

¹*Cocoa: Congestion Control Aware Queuing*, Bachl, Fabini, Zseby, 2020

Deep Learning for Buffering

- Have a separate queue for each flow in a switch/router (fair queuing)
- Define some reward function (like high throughput, low delay)
- Use **Deep Learning**
 - to **fingerprint** each flow
 - and to use the **correct buffer size** for the flow **based on experience**, which **maximizes the reward**

Reward

$$\text{Reward} = \text{bandwidth} - \alpha \times \text{queue size}$$

Input features

- queue size
- standard deviation of the queue size
- maximum allowed buffer size
- rate of incoming data
- rate of outgoing data
- time since the last packet loss

Input features 2

- Don't use these features themselves but exponentially weighted moving averages
- Advantage: No need to keep data around except for features themselves

Neural Network

- It is a regression problem: Predict optimal buffer size for the flow based on the inputs
- Fully connected neural network with three layers and leaky ReLU
- Outputs optimal maximum buffer size each time packet is received
 - To save compute it is possible, for example, only output optimum buffer size every 100 ms

Implementation

- ns-3 for network simulation
- Pytorch's C++ API for Deep Learning
- Integrate Pytorch into ns-3
- Develop queuing discipline based on fair queuing
- Each flow is managed by Reinforcement Learning based on Deep Learning

Offline Learning

- Randomly sample bandwidth, delay, congestion control and flow duration
- For each flow, at random time, launch experiment
- Experiment is A/B testing:
 - Continue with current buffer size +1 and current buffer size -1 (two different experiment)
 - Let the neural network learn buffer size that performed better regarding reward
- Train using L1 loss (mean absolute error), using a couple of results together as a batch


Training takes several hundred thousand flows to converge.

Visualizing success


Should learn

- Larger bandwidth \rightarrow larger buffer
- Larger delay \rightarrow larger buffer


New Reno, varying bandwidth




New Reno, varying delay



Bic, varying bandwidth



Bic, varying delay



Correlations

Correlations calculated of each of the above plots:
Correlation of 100% \rightarrow Our Reinforcement Learning system learned successfully

Table: Correlation between bandwidth/delay of the link and output buffer size.

Congestion control	bandwidth	delay
New Reno	98.3%	99.1%
BIC	80.2%	90.2%


Average queue sizes

Average queue is smaller for New Reno: This is expected since it has a larger multiplicative decrease parameter!


→ New Reno needs a larger buffer on average so that the buffer never becomes empty. Our mechanism learned that!

Congestion control	avg. max. queue length
New Reno	22
BIC	8


25 Mbit/s, 15 ms, New Reno, LFQ



6 Mbit/s, 15 ms, New Reno, LFQ



25 Mbit/s, 15 ms, BIC, LFQ



Correlations


Table: Correlation between bandwidth/delay of the link and output buffer size.

Congestion control	bandwidth	delay
New Reno	99.3%	96.5%
BIC	75.4%	62%

Average queue sizes

Congestion control	avg. max. queue length
New Reno	13
BIC	5

6 Mbit/s, 15 ms, New Reno, LFQ



Online Learning

Same as Offline Learning except:

- Don't do A/B testing
- Instead: Use second neural network that outputs expected reward (*Critic Network*) (trained using L2 loss (mean squared error)).
- Either perform experiment A **or** experiment B (either current buffer size -1 packet or +1 packet)
- If result was better than what the Critic Network expected, let first neural network (*Actor Network*) learn to output this buffer size in the future.

Training takes longer (several million training flows)

Correlations

Table: Correlation between bandwidth/delay of the link and output buffer size.


Congestion control	bandwidth	delay
New Reno	86.1%	87.7%
BIC	72.9%	73.6%

Average queue sizes

Congestion control	avg. max. queue length
New Reno	14
BIC	8


Example: Our solution

15 Mbit/s bandwidth, 5 ms delay




Example: FqCoDel (default on home routers)

15 Mbit/s bandwidth, 5 ms delay



Example: Fifo (default on Linux)

15 Mbit/s bandwidth, 5 ms delay



Systematic comparison

Table: 400 experiments; delay 5 – 25 ms; bandwidth 5 – 25 Mbit/s;
New Reno and Bic Congestion Control. Results averaged.

	avg. throughp.	queue size	
		max.	avg.
LFQ, offline $\alpha = 0.01$	13.4	23.9	7.7
LFQ, offline $\alpha = 10$	12.5	12.7	3.4
LFQ, online $\alpha = 10$	12.8	16.1	4.5
FqCoDel	13.7	155.4	15.4
fq 100	11.7	100	51.1
fq 1000	11.9	1000	630.4

Observations

- Scaling of features very important
- Pytorch better to integrate with other code (ns-3) than TensorFlow
- We thought simulating a flow in ns-3 would be faster than running it in the real world. It is not.

Conclusion

- LFQ is based on fair queuing
- It fingerprints each flow
- It learns to optimize a reward function
- It achieves high throughput and low delay when compared to competing solutions
- It has low computational overhead
- We envision deployment close to end users (on switches, routers)

LFQ: Online Learning of Per-flow Queuing Policies using Deep Reinforcement Learning

Maximilian Bachl, maximilian.bachl@tuwien.ac.at
Joachim Fabini
Tanja Zseby

Technische Universität Wien, Vienna, Austria